Programmieren mit R: Seminararbeit 2

Daniyar Akhmetov (5127348)
Marcelo Rainho Avila (4679876)
Xuan Son Le (4669361)

Abgabedatum: 19.12.2017

Contents

1 Part I: *Functions* (15 points)

1.1 Functions I:

Define a function which given an atomic vector x as argument, returns x after removing missing values

```
dropNa <- function(x) {
  # takes an atomic vector as an argument and returns it without missing values
  #
  # Args:
  #   x: a numeric vector
  #
  # Returns:
  #   atomic vector without NA values
  x[!is.na(x)]
}
```

```
all.equal(dropNa(c(1, 2, 3, NA, 1, 2, 3)), c(1, 2, 3, 1, 2, 3))
```

```
## [1] TRUE
```

1.2 Functions II:

Part I

Write a function meanVarSdSe that takes a numeric vector x as argument. The function should return a named numeric vector that contains the mean, the variance, the standard deviation and the standard error of x.

```
meanVarSdSe <- function(x){
  # computes mean, variance, standard deviation and standard error
  #
  # Args:
  #   x: a numeric vector
  #
  # Returns:
  #   mean, variance, standard deviation and standard error of input vector
  c(mean = mean(x),
    var = var(x),
    sd = sd(x),
    se = sd(x) / sqrt(length(x))
  )
}
```

```
# test
x <- 1:100
meanVarSdSe(x)
```

```
##       mean        var         sd         se
##  50.500000 841.666667  29.011492   2.901149
```

Part II

Look at the following code sequence. What result do you expect?

```r
x <- c(NA, 1:100)
meanVarSdSe(x)
```

The code returns NA values for each statistic computed, which is the output of each function when using the default (FALSE) argument for removing NA's (`na.rm.`)

```r
meanVarSdSe <- function(x, ...){
  # computes mean, variance, standard deviation and standard error
  #
  # Args:
  #   x: a numeric vector
  #
  # Returns:
  #   mean, variance, standard deviation and standard error of input vector
  c(mean = mean(x, ...),
    var = var(x, ...),
    sd = sd(x, ...),
    se = sd(x, ...) / sqrt(length(which(!is.na(x))))
  )
}


# test
meanVarSdSe(x, na.rm = TRUE)
```

```
##       mean        var         sd         se
##  50.500000 841.666667  29.011492   2.901149
```

Part III

Write an alternative version of meanVarSdSe in which you make use of the function definition **dropNa** from the above exercise.

```r
meanVarSdSe <- function(x, dropMissing = TRUE){
  # computs mean, variance, standard deviation and standard error while using
  # dropNa function per default
  #
  # Args:
  #   x: a numeric vector
  #
  # Returns:
  #   mean, variance, standard deviation and standard error of input vector
  if (dropMissing) {
    x <- dropNa(x)
  }
  c(mean = mean(x),
    var = var(x),
    sd = sd(x),
    se = sd(x) / sqrt(length(x))
  )
}


# test
meanVarSdSe(c(x, NA))
```

```
##       mean        var         sd         se
```

```
##  50.500000 841.666667   29.011492    2.901149
```

1.3   Functions III:

Write an infix function %or% that behaves like the logical operator |

```r
"%or%" <- function(x,y) {
  # logical operator OR:
  # TRUE OR TRUE = TRUE
  # TRUE OR FALSE = TRUE
  # FALSE OR TRUE = TRUE
  # FALSE OR FALSE = FALSE

  ifelse(x == TRUE, TRUE,
         ifelse(y == TRUE, TRUE, FALSE))
}

# test
c(TRUE, FALSE, TRUE, FALSE) %or% c(TRUE, TRUE, FALSE, FALSE)
```

```
## [1]  TRUE  TRUE  TRUE FALSE
```

2   Part II: *Scoping and related topics* (15 points)

Scoping I

Explain the results of the three function calls

```r
x <- 5
y <- 7
f <- function() x * y
g <- function(x = 2, y = x) x * y
```

```r
f()      # call 1
```

```
## [1] 35
```

Function `f()` does not require any arguments and is defined as the product of `x` and `y`. It raises an error if `x` or `y` are not are not defined in the global (or any other parent) environment.

```r
g()      # call 2
```

```
## [1] 4
```

Function `g()` takes two arguments, which are `x` with a default value of 2 and `y`, which per default is assigned to the value of `x`. The function returns the product of `x` and `y`. In call 2, function `g()` is called without any arguments. In this case, x is set to 2 and y is equal to x. At first glance, it is unclear if R's interpreter will take the globally assigned value for x (here: 5) or if should take the local variable `x`, which is 2 per default. It turns out that the default arguments are evaluated in the local environment (that is, *inside* the actual function). Therefore, when calling `g()` with default arguments, y is set to equal the *local x* value, which is set to the default value of 2.

```r
g(y=x)   # call 3
```

```
## [1] 10
```

In *call 3*, the argument for `x` is not omitted, so `x` will get the default value 2. Further, `y` is assigned to `x`. Differently to *call 2*, the value of `y` is evaluated when calling the function. Thus `y` is explicitly assigned to the *global x* (in this case 5).

Scoping II

Why and how does the following code work?

```r
t <- matrix(1:6, ncol = 3, byrow = TRUE)
t(t)
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

Here the variable `t` is defined to be a 2 by 3 matrix in the *Global Environment*. However, `t()` is also a function from base R for computing the transpose of a matrix or data.frame. When calling `t(t)`, the parentheses indicate that R should first look for a function `t()` and skip non-function objects, and then apply this function to a object called `t`. R searches for a function `t()` and finds it in the `package:base` environment and calculates the transpose of the previously defined matrix `t`. The call `t(t)` only works because these objects with the same name are not created within the same environment, even when taking into consideration that one is a function and the other is a matrix object.

Scoping III

Why do the results of t(T) and t(t) differ?

```r
t <- function(...) matrix(...)
T <- t(1:6, ncol = 3, byrow = TRUE)
t(T)
```

```
##      [,1]
## [1,]    1
## [2,]    4
## [3,]    2
## [4,]    5
## [5,]    3
## [6,]    6
```

In this scenarion, `t()` is a function defined in the global environment that takes *any* arguments and passes them onto the base R `matrix()` function, which creates a matrix from a given set of values. Further, `T` is a 2 by 3 matrix. However, applying `t()` to `T` is the same as calling `matrix(T)`. This returns, somewhat surprisingly, a 6 by 1 matrix. This dimensionality distortion is due to the fact that a matrix object in R is, under the hood, a *long* one-dimensional atomic vector with a `dim` attribute indicating the number of rows and columns. The default value for the number of rows and columns in `matrix()` function is 1. Therefore, when calling `matrix()` on a matrix object (and not defining a different number of rowns or columns), the matrix `T` gets "unwrapped" into the underlying one-dimensional row vector (or a six-dimensional column vector).

The ordering [1,4,2,5,3,6] rather than [1,2,..,6] is due to `byrow = TRUE` argument when constructing the matrix. The same matrix could, for instance, be created by changing the dimensions of an atomic vector, as following,

```r
aMatrix <- c(1,4,2,5,3,6)
dim(aMatrix) <- c(2,3)
```

which exemplifies the underlying nature of a matrix object in R.

```r
t <- function(...) matrix(...)
t <- t(1:6, ncol = 3, byrow = TRUE)
t(t)
```

```
##      [,1] [,2]
## [1,]    1    4
## [2,]    2    5
## [3,]    3    6
```

In the above scenario, the first line basically maps the base function `matrix()` to `t()`. In the second line, `t()`, while still mapped as a function, creates a 2 by 3 matrix, which is also mapped to `t`, overwriting the previous definition, since we cannot have multiple objects with the same name within the same environment. In the third line. When calling `t(t)`, R searches for the **function `t()`** and finds the transpose function from base R and not the already overwritten function in the Global Environment.

Dynamic Lookup

Explain the results of the five function calls and why the rm function in line 1 is important.

```r
rm(list = ls(all.names = TRUE))
f <- function(x, y = x + 1) x + y
x <- 3
f(2) # call 1
## [1] 5
x <- 5
f(2) # call 2
## [1] 5

f <- function(y = x + 1) x + y
x <- 3
f(2) # call 3
## [1] 5
x <- 5
f(2) # call 4
## [1] 7
f()  # call 5
## [1] 11
```

Call 1 and 2 This function requires one argument, `x`, but also accepts a second argument, `y`. It returns the sum of `x` and `y`, with `y` equals `x + 1` per default.
This exemplifies the lazy evaluation property in R. By the time the function is created `x` doesn't exist and only when the function `f(2)` is called, 2 is passed as an argument for `x` and and `y` gets assigned to the expression `x + 1` from the newly created function environment.) The global value of `x` does not affect the function. Thus, call 1 and 2 returns the same result. One can affirm that `f()` is (weakly) self contained. That means, the values from global or parent environments don't affect the output of the function

Call 3 and 4 This function accepts only one argument, `y`, but a second variable, `x`, is required for it to work. Since this variable is not created within the function environment, R will *go up* the search path looking for a variable called `x`. Even when passing the same argument for `y`, different values of `x` will yield different results, which can be seen on the results of call 3 and call 4. This function is, therefore, not self-contained.

Call 5 Since the default value for `y` is `x + 1` and no argument is passed in the last call, the function returns the value for `x + (x + 1)`, where `x` is the defined in the global environment.

rm function The rm function is useful to make sure that previously assigned values do not generate unexpected results. Nevertheless, it is probably better to make sure this line is not necessary, reducing the risks of getting one's computer burned[1].

---

[1]see discussion on https://twitter.com/hadleywickham/status/940021008764846080